

# Customer Churn Prediction Project

The objective of this project is to develop a machine learning model that accurately predicts monthly user churn for Taxiride navigation app.

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```

```
%matplotlib inline
```

```
df = pd.read_csv('dataset.csv')
df.head(5)
```

	ID	label	sessions	drives	total_sessions
n_days_after_onboarding \					
0	0	retained	283	226	296.748273
2276					
1	1	retained	133	107	326.896596
1225					
2	2	retained	114	95	135.522926
2651					
3	3	retained	49	40	67.589221
15					
4	4	retained	84	68	168.247020
1562					

	total_navigations_fav1	total_navigations_fav2	driven_km_drives	\
0	208	0	2628.845068	
1	19	64	13715.920550	
2	0	0	3059.148818	
3	322	7	913.591123	
4	166	5	3950.202008	

	duration_minutes_drives	activity_days	driving_days	device
0	1985.775061	28	19	Android
1	3160.472914	13	11	iPhone
2	1610.735904	14	8	Android
3	587.196542	7	3	iPhone
4	1219.555924	27	18	Android

```
df.shape
```

```
(14999, 13)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 14999 entries, 0 to 14998
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	ID	14999 non-null	int64
1	label	14299 non-null	object
2	sessions	14999 non-null	int64
3	drives	14999 non-null	int64
4	total_sessions	14999 non-null	float64
5	n_days_after_onboarding	14999 non-null	int64
6	total_navigations_fav1	14999 non-null	int64
7	total_navigations_fav2	14999 non-null	int64
8	driven_km_drives	14999 non-null	float64
9	duration_minutes_drives	14999 non-null	float64
10	activity_days	14999 non-null	int64
11	driving_days	14999 non-null	int64
12	device	14999 non-null	object

```
dtypes: float64(3), int64(8), object(2)
```

```
memory usage: 1.5+ MB
```

```
df.isnull().sum()
```

ID	0
label	700
sessions	0
drives	0
total_sessions	0
n_days_after_onboarding	0
total_navigations_fav1	0
total_navigations_fav2	0
driven_km_drives	0
duration_minutes_drives	0
activity_days	0
driving_days	0
device	0

```
dtype: int64
```


```
df = df.dropna(subset=['label'])
```

```
df.isnull().sum()
```


ID	0
label	0
sessions	0
drives	0
total_sessions	0
n_days_after_onboarding	0
total_navigations_fav1	0

```
total_navigations_fav2    0
driven_km_drives          0
duration_minutes_drives    0
activity_days              0
driving_days               0
device                    0
dtype: int64

sns.countplot(x='label',data = df)
<Axes: xlabel='label', ylabel='count'>
```



```
sns.countplot(x='device',data = df)
<Axes: xlabel='device', ylabel='count'>
```



```
df.columns
Index(['ID', 'label', 'sessions', 'drives', 'total_sessions',
      'n_days_after_onboarding', 'total_navigations_fav1',
      'total_navigations_fav2', 'driven_km_drives',
      'duration_minutes_drives',
      'activity_days', 'driving_days', 'device'],
      dtype='object')

num_col = df.select_dtypes(include=['number']).columns.tolist()
print(num_col)

['ID', 'sessions', 'drives', 'total_sessions',
 'n_days_after_onboarding', 'total_navigations_fav1',
 'total_navigations_fav2', 'driven_km_drives',
 'duration_minutes_drives', 'activity_days', 'driving_days']

df.set_index('ID', inplace=True)

num_col

['ID',
 'sessions',
 'drives',
 'total_sessions',
```

```

'n_days_after_onboarding',
'total_navigations_fav1',
'total_navigations_fav2',
'driven_km_drives',
'duration_minutes_drives',
'activity_days',
'driving_days']

num_col.remove('ID')

num_col


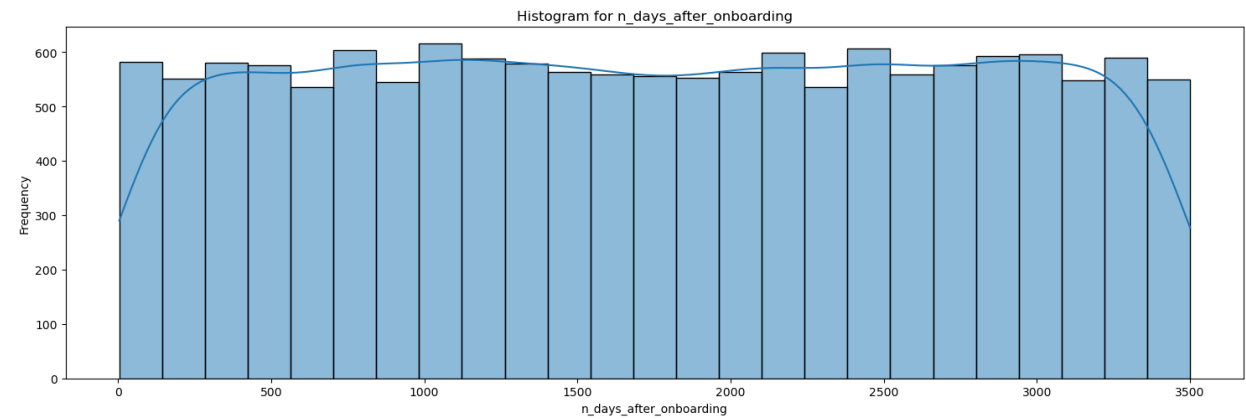
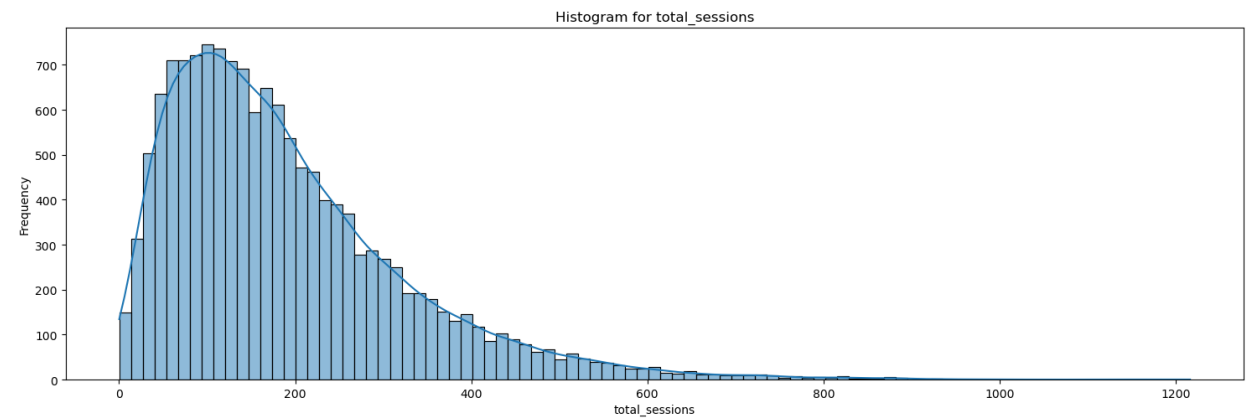
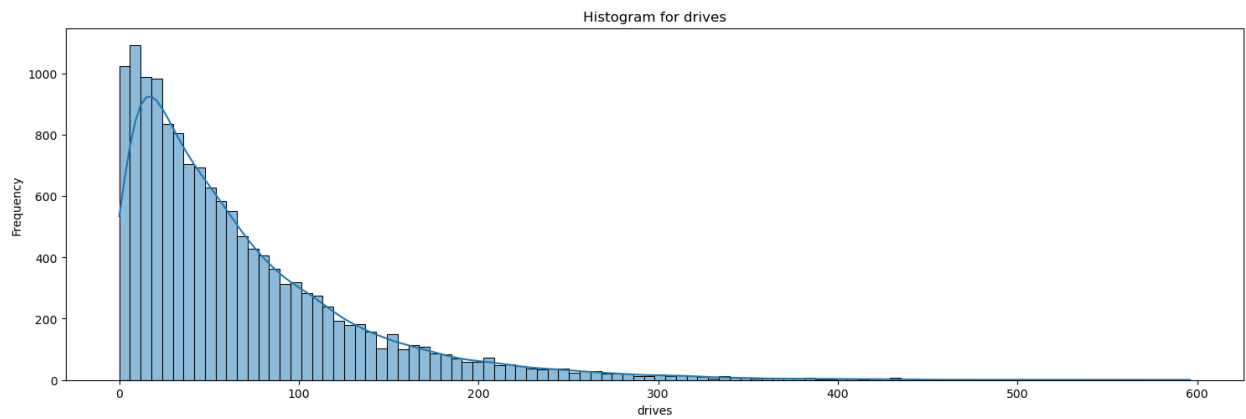
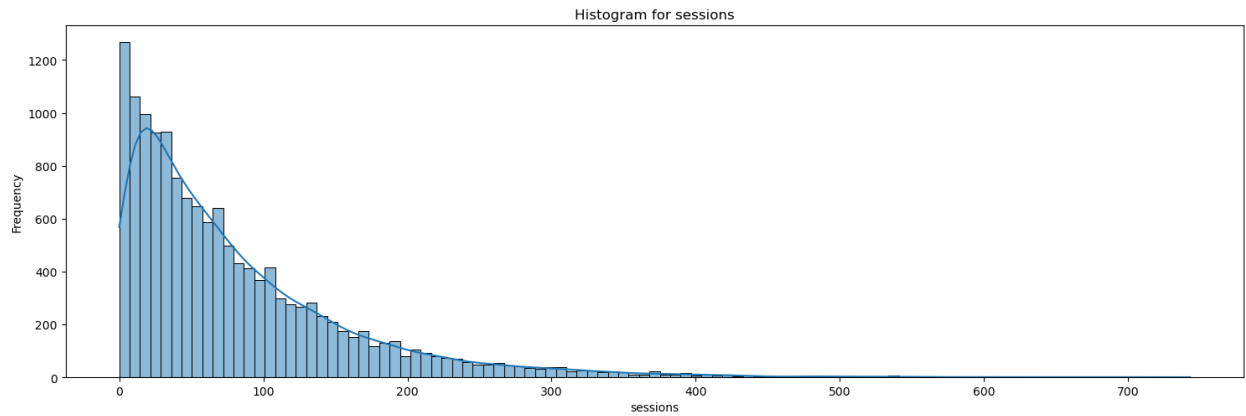
['sessions',
'drives',
'total_sessions',
'n_days_after_onboarding',
'total_navigations_fav1',
'total_navigations_fav2',
'driven_km_drives',
'duration_minutes_drives',
'activity_days',
'driving_days']

plt.figure(figsize=(15,5*len(num_col)))
plot_index = 1

for column in num_col:
    plt.subplot(len(num_col),1,plot_index)
    sns.histplot(data=df, x=column, kde=True)
    plt.title(f'Histogram for {column}')
    plt.xlabel(column)
    plt.ylabel('Frequency')
    plot_index += 1

plt.tight_layout()
plt.show()

```



```

from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
cat_col = ['label', 'device']
df[cat_col] = df[cat_col].apply(label_encoder.fit_transform)
df.head()

```

	label	sessions	drives	total_sessions	
n_days_after_onboarding \					
ID					
0	1	283	226	296.748273	2276
1	1	133	107	326.896596	1225
2	1	114	95	135.522926	2651
3	1	49	40	67.589221	15
4	1	84	68	168.247020	1562


	total_navigations_fav1	total_navigations_fav2	
driven_km_drives \			
ID			
0	208	0	2628.845068
1	19	64	13715.920550
2	0	0	3059.148818
3	322	7	913.591123
4	166	5	3950.202008

	duration_minutes_drives	activity_days	driving_days	device
ID				
0	1985.775061	28	19	0
1	3160.472914	13	11	1
2	1610.735904	14	8	0
3	587.196542	7	3	1
4	1219.555924	27	18	0

```

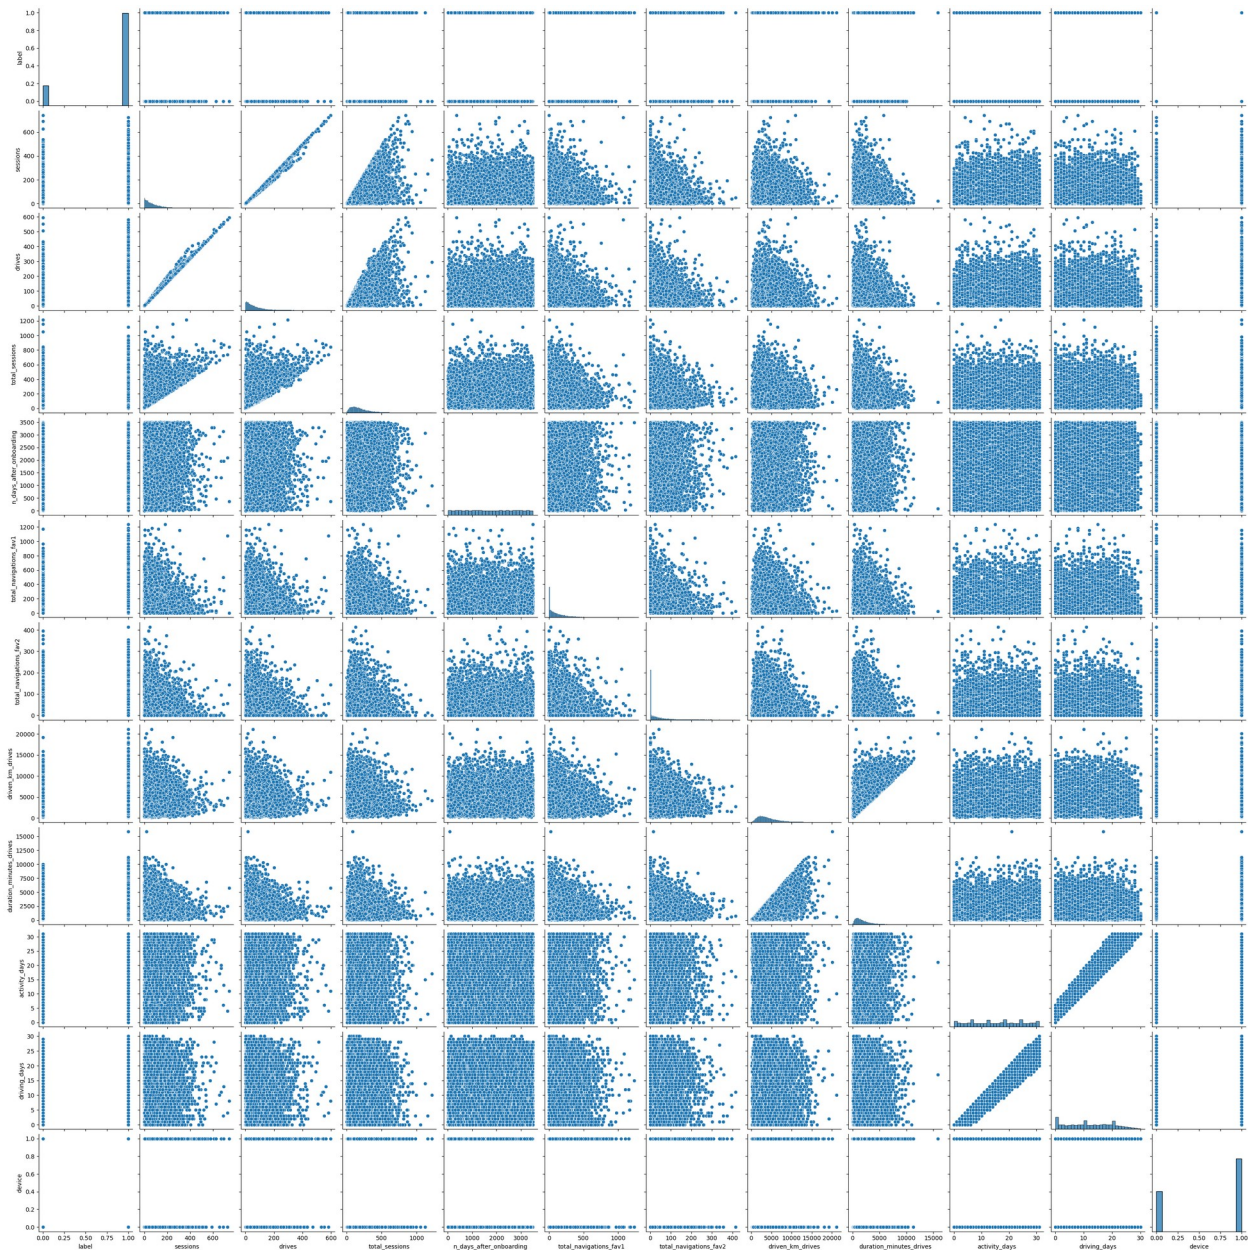
corr_matrix = df.corr()
plt.figure(figsize=(12,10))
sns.heatmap(corr_matrix, annot=True, cmap='cividis', fmt='.2f')
plt.title('Correlation Heatmap before Standardization')
plt.show()

```




```
sns.pairplot(df)  
plt.show()
```





```
corr_spearman = df.corr(method='spearman')
plt.figure(figsize=(12,10))
sns.heatmap(corr_spearman, annot=True, cmap='cividis',
fmt='.2f', cbar=True, square=True)
plt.title('Correlation Heatmap Spearman method')
plt.show()
```



```
# Applying scaling technique for diff numerical columns
# log transformation = sessions,drives,total_sessions, fav1, fav2,
# standardtransformation = activity_days, driving_days,
# n_days_after_onboarding

# after that lasso will be applied to figure out less imp feature
from sklearn.preprocessing import RobustScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split

X = df.drop(columns=['label'])
y = df['label']

X_train,X_test,y_train,y_test = train_test_split(X, y, test_size=0.3,
```

```

random_state=1)

columns_for_robust_scale =
['activity_days', 'driving_days', 'n_days_after_onboarding']
columns_for_log =
['sessions', 'drives', 'total_sessions', 'total_navigations_fav1', 'total_
navigations_fav2', 'driven_km_drives', 'duration_minutes_drives']

robust_scaler = RobustScaler()
X_train[columns_for_robust_scale]=robust_scaler.fit_transform(X_train[
columns_for_robust_scale])
X_test[columns_for_robust_scale]=robust_scaler.transform(X_test[column
s_for_robust_scale])

X_train[columns_for_log]= np.log1p(X_train[columns_for_log])
X_test[columns_for_log]= np.log1p(X_test[columns_for_log])

print('Transformed X_train : \n', X_train.head())
print('\nTransformed X_test : \n', X_test.head())

```

Transformed X\_train :

	sessions	drives	total_sessions	n_days_after_onboarding	\
ID					
13832	4.912655	4.736198	4.916879	0.703259	
7463	5.187386	4.962845	5.370140	-0.579760	
14967	5.433722	5.241747	6.334882	-0.201830	
588	4.276666	4.143135	5.166819	0.486564	
8297	4.709530	4.532599	5.286255	0.785020	

	total_navigations_fav1	total_navigations_fav2
driven_km_drives \		
ID		

13832	6.102559	4.356709
8.704022		
7463	3.970292	0.000000
8.432454		
14967	1.945910	4.770685
7.938738		
588	3.583519	3.555348
9.140724		
8297	0.000000	0.000000
8.722623		

	duration_minutes_drives	activity_days	driving_days	device
ID				
13832	8.024135	-0.875	-0.714286	1
7463	7.381038	0.875	1.000000	1

14967	6.588099	0.875	0.714286	1
588	8.797199	0.125	0.214286	1
8297	7.832262	-0.500	-0.285714	1

Transformed X\_train :

	sessions	drives	total_sessions	n_days_after_onboarding	\
ID					
7475	4.564348	4.343805	5.645627	0.118353	
7387	3.258097	3.258097	4.770404	0.280732	
12121	3.218876	3.218876	5.066297	0.788451	
11451	2.944439	2.833213	4.913304	-0.296169	
14149	4.219508	4.219508	5.841180	-0.204117	

	total_navigations_fav1	total_navigations_fav2
driven_km_drives \		
ID		
7475	6.373320	5.370638
8.552705		
7387	5.823046	2.708050
7.659630		
12121	5.198497	0.000000
8.232976		
11451	4.709530	4.488636
8.445051		
14149	4.941642	0.000000
7.930120		

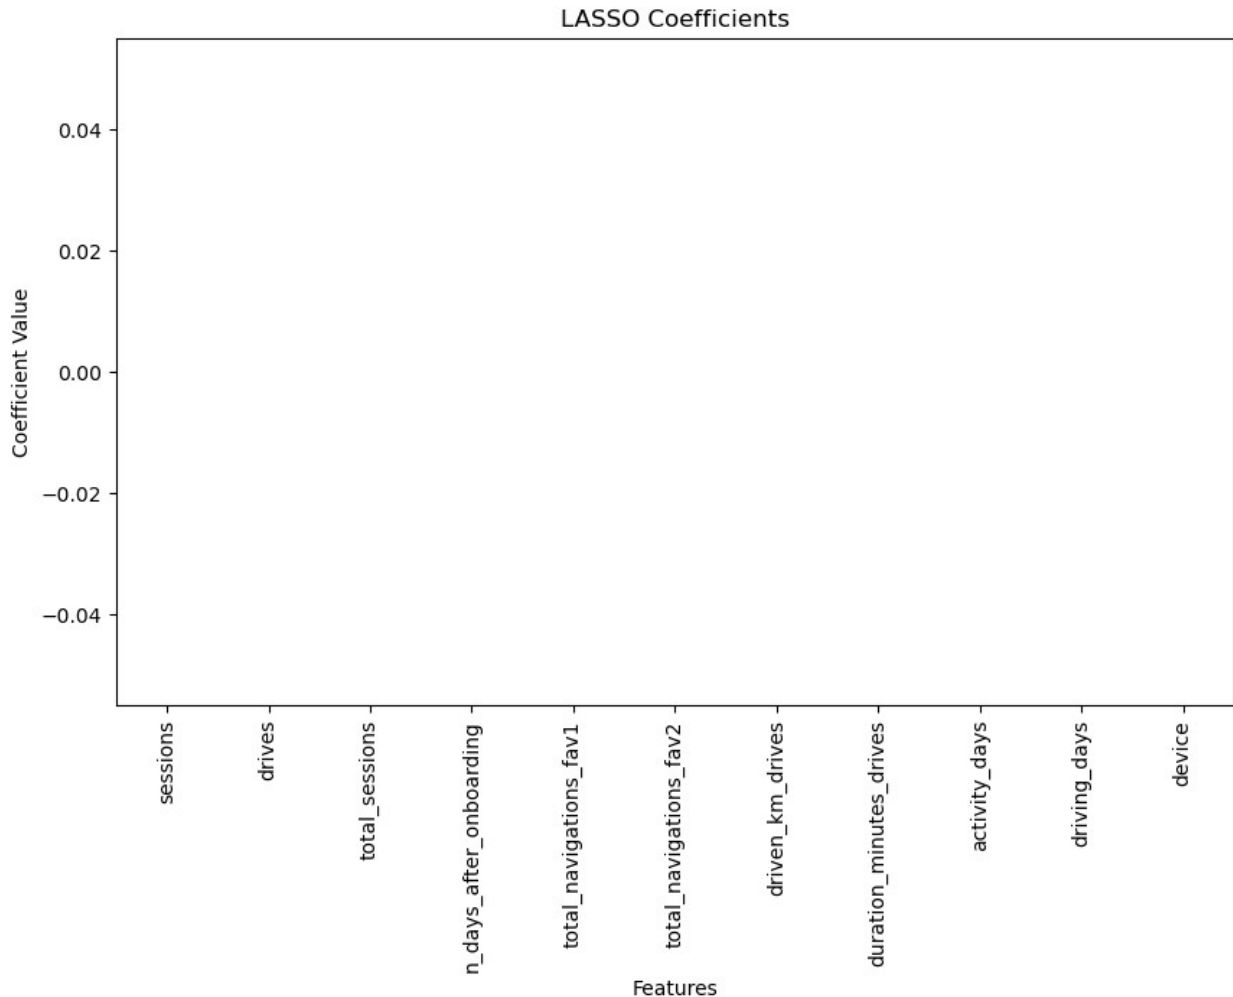
	duration_minutes_drives	activity_days	driving_days	device
ID				
7475	6.787305	0.1250	-0.071429	1
7387	6.866728	-0.6250	-0.714286	1
12121	7.685811	-0.3125	-0.071429	1
11451	8.097410	-0.7500	-0.857143	1
14149	6.875838	-0.3750	-0.142857	0

```
from sklearn.linear_model import Lasso
```

```
lasso = Lasso(alpha=0.1)
lasso.fit(X_train,y_train)
```

```
lasso_coef = lasso.coef_
feature_importance = pd.Series(lasso_coef, index=X.columns)
```

```
plt.figure(figsize=(10, 6))
feature_importance.plot(kind='bar')
plt.title('LASSO Coefficients')
plt.xlabel('Features')
plt.ylabel('Coefficient Value')
plt.show()
```



```
selected_features = feature_importance[feature_importance != 0]
print("Selected Features:\n", selected_features)
```

```
Selected Features:
Series([], dtype: float64)
```

```
X_train =
X_train.drop(columns=['sessions', 'total_navigations_fav1', 'total_navigations_fav2', 'driven_km_drives', 'device', 'activity_days'])
X_test =
X_test.drop(columns=['sessions', 'total_navigations_fav1', 'total_navigations_fav2', 'driven_km_drives', 'device', 'activity_days'])
```

```
X_train.head()
```

	drives	total_sessions	n_days_after_onboarding	\
ID				
13832	4.736198	4.916879		0.703259
7463	4.962845	5.370140		-0.579760
14967	5.241747	6.334882		-0.201830

588	4.143135	5.166819	0.486564
8297	4.532599	5.286255	0.785020

	duration_minutes_drives	driving_days
ID		
13832	8.024135	-0.714286
7463	7.381038	1.000000
14967	6.588099	0.714286
588	8.797199	0.214286
8297	7.832262	-0.285714

```

from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import cross_val_score, GridSearchCV
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix, f1_score, precision_score, recall_score

```

## Model Training

```

!pip install xgboost
from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier
models = {
    'Logistic Regression': LogisticRegression(max_iter=200),
    'K-Nearest Neighbors': KNeighborsClassifier(n_neighbors=5),
    'Random Forest': RandomForestClassifier(n_estimators=100),
    'Support Vector Machine': SVC(),
    'Gradient Boosting': GradientBoostingClassifier(),
    'XGBoost': XGBClassifier()
}

results = {}

for model_name, model in models.items():
    cv_scores = cross_val_score(model, X_train, y_train, cv=5) # 5-
fold cross-validation
    results[model_name] = {
        'CV Scores': cv_scores,
        'Mean CV Score': np.mean(cv_scores)
    }

for model_name, result in results.items():
    print(f"{model_name}:\n Cross-Validation Scores: {result['CV
Scores']}\n Mean CV Score: {result['Mean CV Score']:.4f}\n")

```

Requirement already satisfied: xgboost in c:\users\bansa\anaconda3\lib\site-packages (2.1.2)



```
Requirement already satisfied: numpy in c:\users\bansa\anaconda3\lib\
site-packages (from xgboost) (1.26.4)
Requirement already satisfied: scipy in c:\users\bansa\anaconda3\lib\
site-packages (from xgboost) (1.11.4)
Logistic Regression:
  Cross-Validation Scores: [0.82267732 0.82817183 0.82967033
0.82317682 0.82808596]
  Mean CV Score: 0.8264

K-Nearest Neighbors:
  Cross-Validation Scores: [0.81618382 0.8006993 0.7962038
0.81118881 0.8045977 ]
  Mean CV Score: 0.8058

Random Forest:
  Cross-Validation Scores: [0.82117882 0.81918082 0.82217782
0.82567433 0.82008996]
  Mean CV Score: 0.8217

Support Vector Machine:
  Cross-Validation Scores: [0.82417582 0.82417582 0.82417582
0.82467532 0.82458771]
  Mean CV Score: 0.8244

Gradient Boosting:
  Cross-Validation Scores: [0.82417582 0.82717283 0.82817183
0.83066933 0.82058971]
  Mean CV Score: 0.8262

XGBoost:
  Cross-Validation Scores: [0.80669331 0.8036963 0.80619381
0.81868132 0.81209395]
  Mean CV Score: 0.8095
```

Logistic Regression and Gradient Boosting classifiers have high CV score among others.

```
## Hyper parameter tuning of Logistic regression and Gradient Boosting

Logistic_param_grid = {
    'penalty': ['l1', 'l2', 'elasticnet', 'none'],
    'C': [0.01, 0.1, 1.0, 10, 100],
    'solver': ['liblinear', 'lbfgs', 'saga'],
    'class_weight': [None, 'balanced']
}
```

```

logistic_model = LogisticRegression()

logi_grid_search = GridSearchCV(estimator=logistic_model,
param_grid=Logistic_param_grid, cv=5, scoring='accuracy', n_jobs=-1 )

logi_grid_search.fit(X_train, y_train)

print("Best Parameters:", logi_grid_search.best_params_)
print("Best CV Score:", logi_grid_search.best_score_)

Best Parameters: {'C': 1.0, 'class_weight': None, 'penalty': 'l2',
'solver': 'liblinear'}
Best CV Score: 0.8271558027180216

final_logistic_model = logi_grid_search.best_estimator_
logi_y_pred = final_logistic_model.predict(X_test)

logi_accuracy = accuracy_score(y_test, logi_y_pred)

print("Test Set Accuracy:", logi_accuracy)
print("Classification Report:\n", classification_report(y_test,
logi_y_pred))

```

Test Set Accuracy: 0.8200466200466201

Classification Report:

	precision	recall	f1-score	support
0	0.54	0.05	0.09	778
1	0.82	0.99	0.90	3512
accuracy			0.82	4290
macro avg	0.68	0.52	0.50	4290
weighted avg	0.77	0.82	0.75	4290